

Introduction to sponge-based cryptography

Part 1: KECCAK and SHA-3

Joan DAEMEN

STMicroelectronics and Radboud University

SPACE 2016

Hyderabad, India, December 14, 2016

Outline


- 1 The SHA-3 competition
- 2 The sponge construction
- 3 Inside KECCAK
- 4 The SHA-3 FIPS
- 5 KangarooTwelve

Outline

- 1 The SHA-3 competition
- 2 The sponge construction
- 3 Inside KECCAK
- 4 The SHA-3 FIPS
- 5 KangarooTwelve

Cryptographic hash functions

- Function h from \mathbf{Z}_2^* to \mathbf{Z}_2^n
- Typical values for n : 128, 160, 256, 512




- Pre-image resistant: it shall take 2^n effort to
 - given y , find x such that $h(x) = y$
- 2nd pre-image resistance: it shall take 2^n effort to
 - given M and $h(M)$, find another M' with $h(M') = h(M)$
- Collision resistance: it shall take $2^{n/2}$ effort to
 - find $x_1 \neq x_2$ such that $h(x_1) = h(x_2)$
- More general: *should behave like a random oracle*

Ideal function: random oracle \mathcal{RO}

- A random oracle [Bellare-Rogaway 1993] maps:
 - message of variable length
 - to an infinite output string
- Supports queries of following type: (M, ℓ)
 - M : message
 - ℓ : requested number of output bits
- Response Z
 - string of ℓ bits
 - independently and uniformly distributed bits
 - self-consistent: equal M give matching outputs
- The ultimate hash function

Classical way to build hash functions

- Mode of use of a compression function:
 - Fixed-input-length compression function
 - Merkle-Damgård iterating mode



- Property-preserving paradigm
 - hash function *inherits* properties of compression function
 - ...actually block cipher with feed-forward (Davies-Meyer)
- Compression function built on arithmetic-rotation-XOR: **ARX**
- Instances: MD5, SHA-1, SHA-2 (224, 256, 384, 512) ...

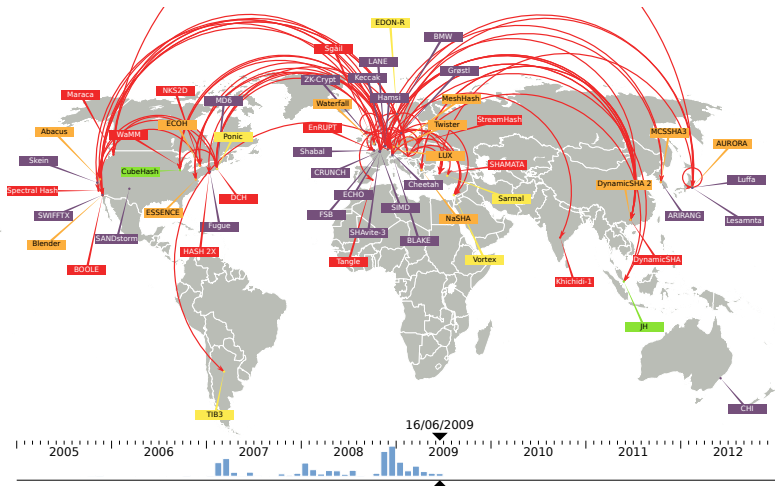
The origins of the SHA-3 competition

- 2005-2006: NIST thinks about having a SHA-3 contest
 - MD5 and standard SHA-1 were damaged by attacks
 - SHA-2 based on the same principles than MD5 and SHA-1
 - open call for SHA-2 successor
 - ...and for analysis, comparisons, etc.
- October 2008: Deadline for proposals
 - *more efficient than SHA-2*
 - output lengths: 224, 256, 384, 512 bits
 - security: collision and (2nd) pre-image resistant
 - specs, reference and optimized code, test vectors
 - design rationale and preliminary analysis
 - patent waiver

The SHA-3 competition

- First round: October 2008 to summer 2009
 - 64 submissions, 51 accepted
 - 37 presented at 1st SHA-3 candidate conf. in Leuven, February 2009
 - many broken by cryptanalysis
 - NIST narrowed down to 14 semi-finalists
- Second round: summer 2009 to autumn 2010
 - analysis presented at 2nd SHA-3 conf. in Santa Barbara, August 2010
 - NIST narrowed down to 5 finalists
- Third round: autumn 2010 to October 2012
 - analysis presented at 3rd SHA-3 conf. in Washington, March 2012
- October 2, 2012: NIST announces KECCAK as SHA-3 winner

NIST SHA-3: the battlefield



[courtesy of Christophe De Cannière]


Outline

- 1 The SHA-3 competition
- 2 The sponge construction**
- 3 Inside KECCAK
- 4 The SHA-3 FIPS
- 5 KangarooTwelve


Sponge origin: RADIOGATÚN

- Initiative to design hash/stream function (late 2005)
 - rumours about NIST call for hash functions
 - starting point: **fixing PANAMA** [Daemen, Clapp, FSE 1998]
 - with long-time colleagues **Gilles Van Assche** and **Michaël Peeters**
 - and ST Italy colleague **Guido Bertoni** joining in
- RADIOGATÚN [Keccak team, NIST 2nd hash workshop 2006]
 - more conservative than PANAMA
 - **arbitrary output length**
 - **expressing security claim** for arbitrary output length function
- Sponge functions [Keccak team, Ecrypt hash, 2007]
 - **random sponge** instead of random oracle as security goal
 - sponge construction calling random permutation
 - *... closest thing to a random oracle with a finite state ...*

The (extended) KECCAK team




The sponge construction




- Generalizes hash function: *extendable output function* (XOF)
- Calls a b -bit **permutation** f , with $b = r + c$
 - r bits of *rate*
 - c bits of *capacity* (security parameter)

Generic security: indistinguishability



- Success probability of distinguishing between:
 - ideal function: a monolithic random oracle \mathcal{RO}
 - construction $\mathcal{S}[\mathcal{F}]$ calling a random permutation \mathcal{F}
- Adversary \mathcal{D} sends queries (M, ℓ) according to algorithm
- Express $\Pr(\text{success}|\mathcal{D})$ as a function of total cost of queries N
- **Problem: in real world, \mathcal{F} is available to adversary**

Generic security: indistinguishability [Maurer et al. (2004)]



- Applied to hash functions in [Coron et al. (2005)]
 - distinguishing mode-of-use from ideal function (\mathcal{RO})
 - covers adversary with access to permutation \mathcal{F} at left
 - additional interface, covered by a *simulator* at right
- Methodology:
 - build \mathcal{P} that makes left/right distinguishing difficult
 - prove bound for advantage given this simulator \mathcal{P}
 - \mathcal{P} may query \mathcal{RO} for acting \mathcal{S} -consistently: $\mathcal{P}[\mathcal{RO}]$

Generic security of the sponge construction

Concept of *advantage*:

$$\Pr(\text{success}|\mathcal{D}) = \frac{1}{2} + \frac{1}{2}\text{Adv}(\mathcal{D})$$

Theorem (Bound on the \mathcal{RO} -differentiating advantage of sponge)

$$A \leq \frac{N^2}{2^{c+1}}$$

A: differentiating advantage of random sponge from random oracle

N: total data complexity

c: capacity

[Keccak team, Eurocrypt 2008]

Implications of the bound

- Let \mathcal{D} : n -bit output pre-image attack. Success probability:
 - for random oracle: $P_{\text{pre}}(\mathcal{D}|\mathcal{RO}) = q2^{-n}$
 - for random sponge: $P_{\text{pre}}(\mathcal{D}|\mathcal{S}[\mathcal{F}]) = ?$
- A distinguisher \mathcal{D} with $A = P_{\text{pre}}(\mathcal{D}|\mathcal{S}[\mathcal{F}]) - P_{\text{pre}}(\mathcal{D}|\mathcal{RO})$
 - do pre-image attack
 - if success, conclude random sponge and \mathcal{RO} otherwise
- But we have a proven bound $A \leq \frac{N^2}{2^{c+1}}$, so

$$P_{\text{pre}}(\mathcal{D}|\mathcal{S}[\mathcal{F}]) \leq P_{\text{pre}}(\mathcal{D}|\mathcal{RO}) + \frac{N^2}{2^{c+1}}$$

- Can be generalized to any attack
- Note that A is independent of output length n

Implications of the bound (cont'd)

- Informally, random sponge is like random oracle for $N < 2^{c/2}$
- Security strength for output length n :
 - collision-resistance: $\min(c/2, n/2)$
 - first pre-image resistance: $\min(c/2, n)$
 - second pre-image resistance: $\min(c/2, n)$
- Proof assumes f is a **random** permutation
 - provably secure against generic attacks
 - ...but not against attacks that exploit specific properties of f
- No security against multi-stage adversaries

Design approach

Hermetic sponge strategy

- Instantiate a **sponge function**
- Claim a security level of $2^{c/2}$

Our mission

Design permutation f without exploitable properties

How to build a strong permutation

- Like a block cipher
 - Sequence of identical rounds
 - Round consists of sequence of simple step mappings
- ...but not quite
 - No key schedule
 - Round constants instead of round keys
 - Inverse permutation need not be efficient

Outline

- 1 The SHA-3 competition
- 2 The sponge construction
- 3 Inside KECCAK**
- 4 The SHA-3 FIPS
- 5 KangarooTwelve

KECCAK $[r, c]$

- Sponge function using the **permutation** KECCAK- f
 - 7 permutations: $b \in \{25, 50, 100, 200, 400, 800, 1600\}$
... from toy over lightweight to high-speed ...
- SHA-3 instance: $r = 1088$ and $c = 512$
 - permutation width: 1600
 - security strength 256: post-quantum sufficient
- Lightweight instance: $r = 40$ and $c = 160$
 - permutation width: 200
 - security strength 80: same as (initially expected from) SHA-1

See [The KECCAK reference] for more details

KECCAK $[r, c]$

- Sponge function using the **permutation** KECCAK- f
 - 7 permutations: $b \in \{25, 50, 100, 200, 400, 800, 1600\}$
... from toy over lightweight to high-speed ...
- SHA-3 instance: $r = 1088$ and $c = 512$
 - permutation width: 1600
 - security strength 256: post-quantum sufficient
- Lightweight instance: $r = 40$ and $c = 160$
 - permutation width: 200
 - security strength 80: same as (initially expected from) SHA-1


See [The KECCAK reference] for more details

KECCAK $[r, c]$

- Sponge function using the **permutation** KECCAK- f
 - 7 permutations: $b \in \{25, 50, 100, 200, 400, 800, 1600\}$
... from toy over lightweight to high-speed ...
- SHA-3 instance: $r = 1088$ and $c = 512$
 - permutation width: 1600
 - security strength 256: post-quantum sufficient
- Lightweight instance: $r = 40$ and $c = 160$
 - permutation width: 200
 - security strength 80: same as (initially expected from) SHA-1


See [The KECCAK reference] for more details

The state: an array of $5 \times 5 \times 2^\ell$ bits




- 5×5 **lanes**, each containing 2^ℓ bits (1, 2, 4, 8, 16, 32 or 64)
- (5×5) -bit **slices**, 2^ℓ of them

The state: an array of $5 \times 5 \times 2^\ell$ bits




- 5×5 **lanes**, each containing 2^ℓ bits (1, 2, 4, 8, 16, 32 or 64)
- (5×5) -bit **slices**, 2^ℓ of them

The state: an array of $5 \times 5 \times 2^\ell$ bits




- 5×5 **lanes**, each containing 2^ℓ bits (1, 2, 4, 8, 16, 32 or 64)
- (5×5) -bit **slices**, 2^ℓ of them

The state: an array of $5 \times 5 \times 2^\ell$ bits




- 5×5 **lanes**, each containing 2^ℓ bits (1, 2, 4, 8, 16, 32 or 64)
- (5×5) -bit **slices**, 2^ℓ of them

The state: an array of $5 \times 5 \times 2^\ell$ bits



- 5×5 **lanes**, each containing 2^ℓ bits (1, 2, 4, 8, 16, 32 or 64)
- (5×5) -bit **slices**, 2^ℓ of them

χ , the nonlinear mapping in KECCAK- f



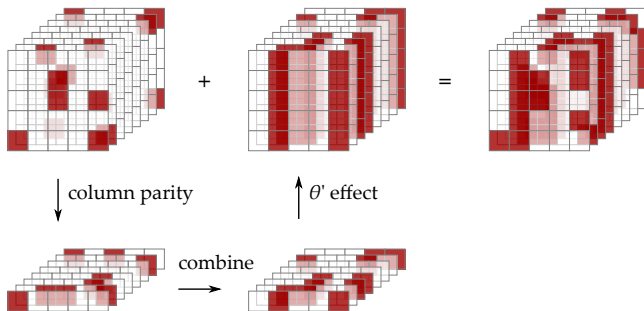
- “Flip bit if neighbors exhibit 01 pattern”
- Operates independently and in parallel on 5-bit rows
- **Cheap**: small number of operations per bit

θ' , a first attempt at mixing bits


- Compute parity $c_{x,z}$ of each column
- Add to each cell parity of neighboring columns:

$$b_{x,y,z} = a_{x,y,z} \oplus c_{x-1,z} \oplus c_{x+1,z}$$

- **Cheap**: two XORs per bit




Diffusion of θ'




$$1 + (1 + y + y^2 + y^3 + y^4) (x + x^4) \\ (\text{mod } \langle 1 + x^5, 1 + y^5, 1 + z^w \rangle)$$

Diffusion of θ' (kernel)



$$1 + (1 + y + y^2 + y^3 + y^4) (x + x^4) \\ (\text{mod } \langle 1 + x^5, 1 + y^5, 1 + z^w \rangle)$$

Diffusion of the inverse of θ'




$$1 + (1 + y + y^2 + y^3 + y^4) (x^2 + x^3) \\ (\text{mod } \langle 1 + x^5, 1 + y^5, 1 + z^w \rangle)$$

ρ for inter-slice dispersion

- We need diffusion between the slices ...
- ρ : cyclic shifts of lanes with offsets

$$i(i+1)/2 \bmod 2^\ell, \text{ with } \begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} 0 & 1 \\ 2 & 3 \end{pmatrix}^{i-1} \begin{pmatrix} 1 \\ 0 \end{pmatrix}$$

- Offsets cycle through all values below 2^ℓ




ι to break symmetry

- XOR of round-dependent constant to lane in origin
- Without ι , the round mapping would be symmetric
 - invariant to translation in the z-direction
 - susceptible to *rotational* cryptanalysis
- Without ι , all rounds would be the same
 - susceptibility to *slide* attacks
 - defective cycle structure
- Without ι , we get simple fixed points (000 and 111)


A first attempt at KECCAK-f

- Round function: $R = \iota \circ \rho \circ \theta' \circ \chi$
- Problem: low-weight periodic trails by chaining:




- χ : propagates unchanged with weight 4
- θ' : propagates unchanged, because all column parities are 0
- ρ : in general moves active bits to different slices ...
...but not always

The Matryoshka property



- Patterns in Q' are z-periodic versions of patterns in Q
- Weight of trail Q' is twice that of trail Q (or 2^n times in general)


π for disturbing horizontal/vertical alignment



$$a_{x,y} \leftarrow a_{x',y'} \text{ with } \begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} 0 & 1 \\ 2 & 3 \end{pmatrix} \begin{pmatrix} x' \\ y' \end{pmatrix}$$

A second attempt at KECCAK- f


- Round function: $R = \iota \circ \pi \circ \rho \circ \theta' \circ \chi$
- Solves problem encountered before:



- π moves bits in same column to different columns!


Almost there, still a final tweak ...

Tweaking θ' to θ



$$1 + (1 + y + y^2 + y^3 + y^4) (x + x^4 z) \\ (\text{mod } \langle 1 + x^5, 1 + y^5, 1 + z^w \rangle)$$

Inverse of θ



$$1 + (1 + y + y^2 + y^3 + y^4) \mathbf{Q},$$

with $\mathbf{Q} = 1 + (1 + x + x^4 z)^{-1} \bmod \langle 1 + x^5, 1 + z^w \rangle$

- \mathbf{Q} is dense, so:
 - Diffusion from single-bit output to input very high
 - Increases resistance against LC/DC and algebraic attacks

KECCAK- f summary

- Round function:

$$R = \iota \circ \chi \circ \pi \circ \rho \circ \theta$$

- Number of rounds: $12 + 2\ell$

- KECCAK- $f[25]$ has 12 rounds
- KECCAK- $f[1600]$ has 24 rounds

- Efficiency [KECCAK implementation overview]

- high level of parallelism
- flexibility: bit-interleaving
- software: fast on wide range of CPU
- dedicated hardware: very fast
- suited for protection against side-channel attack

[Debande, Le and Keccak team, HASP 2012 + ePrint 2013/067]

Crunchy Crypto Collision and Preimage Contest


- Goal:
 - Motivate 3rd-party cryptanalysis
 - Give an instant view on current state-of-the-art
- Scope: 1 to 12 rounds, including smaller instances
 - $\text{KECCAK}[r = 40, c = 160]$,
 - $\text{KECCAK}[r = 240, c = 160]$,
 - $\text{KECCAK}[r = 640, c = 160]$ and
 - $\text{KECCAK}[r = 1440, c = 160]$.
- Results so far:
 - Preimages found for up to 4 rounds
 - Collisions found for up to 5 rounds
 - both by team incl. Meicheng Liu (China) and Jian Guo (Singapore)

http://keccak.noekeon.org/crunchy_contest.html


Performance in hardware: high-speed core

From Kris Gaj's presentation at SHA-3, Washington 2012:

ASIC



Stratix III FPGA



Outline

- 1 The SHA-3 competition
- 2 The sponge construction
- 3 Inside KECCAK
- 4 The SHA-3 FIPS**
- 5 KangarooTwelve

The long road to the SHA-3 FIPS

- February 2013: NIST-KECCAK-team meeting
 - SHA-2 replacement by now less urgent
 - ...but KECCAK is more than just hashing!
- NIST disseminates joint SHA-3 proposal
- Summer 2013: Snowden revelations
 - alleged NSA backdoor in DUAL EC DRBG
 - SHA-3 proposal framed as “NIST weakening KECCAK”
- Early 2014: standard takes shape addressing public concerns
- Friday, April 4, 2014: draft FIPS 202 for public comments
- August 2014: NIST announces plans at SHA-3 conference
- August 2015: FIPS 202 official publication

FIPS 202: what is inside?

■ Content

- KECCAK instances for
 - 4 hash functions
 - 2 XOFs
- KECCAK- f all 7 block widths
 - even reduced-round versions
 - unlike AES FIPS that has only 1 of the 5 Rijndael widths
- sponge construction

■ Concept: toolbox for building other functions

- dedicated *special publications* (NIST SP 800-185) soon to be published
- covers parallel tree hashing, MAC and key derivation

<http://csrc.nist.gov/groups/ST/hash/sha-3/Aug2014/index.html>

XOF: eXtensible Output Function

“XOF: a function in which the output can be extended to any length.”

- Good for full domain hash, stream ciphers and key derivation

[Ray Perlner, SHA 3 workshop 2014]

- Quite natural for sponge
 - keeps state and delivers more output upon request
 - bits of output do not depend on the number of bits requested
- Allows simplification:
 - instead of separate hash functions per output length
 - a single XOF can cover **all** use cases:

$$H_{-256}(M) = \lfloor \text{XOF}(M) \rfloor_{256}$$

Domain separation

- Some protocols and applications need
 - multiple hash functions or XOFs
 - that should be *independent*
- With a single XOF?
- Yes: using **domain separation**
 - output of $\text{XOF}(M||0)$ and $\text{XOF}(M||1)$ are independent
 - ...unless XOF has a cryptographic weakness
- Generalization to 2^n functions with D an n -bit *diversifier*

$$\text{XOF}_D(M) = \text{XOF}(M||D)$$

- Variable-length diversifiers: suffix-free set of strings


The XOFs and hash functions in FIPS 202

- Four drop-in replacements identical to those in KECCAK submission
- Two *extendable output functions* (XOF)
- Tree-hashing ready: **SAKURA** coding [Keccak team, ePrint 2013/231]

XOF	SHA-2 drop-in replacements
KECCAK[c = 256](M 11 11)	
	$\lfloor \text{KECCAK}[c = 448](M 01) \rfloor_{224}$
KECCAK[c = 512](M 11 11)	
	$\lfloor \text{KECCAK}[c = 512](M 01) \rfloor_{256}$
	$\lfloor \text{KECCAK}[c = 768](M 01) \rfloor_{384}$
	$\lfloor \text{KECCAK}[c = 1024](M 01) \rfloor_{512}$
SHAKE128 and SHAKE256	SHA3-224 to SHA3-512

All 6 functions in 9 tweets: <https://twitter.com/tweetfips202>

Tree hashing



Features:

- hash recomputation when modifying small part of file
- parallelizable
- performance:

function	instruction	cycles/byte
$\text{KECCAK}[c = 256] \times 1$	x86_64	7.70
$\text{KECCAK}[c = 256] \times 2$	AVX2 (128-bit only)	5.30
$\text{KECCAK}[c = 256] \times 4$	AVX2	2.87

CPU: Haswell with AVX2 256-bit SIMD

Outline

- 1 The SHA-3 competition
- 2 The sponge construction
- 3 Inside KECCAK
- 4 The SHA-3 FIPS
- 5 KangarooTwelve**

Safety margin: from *rock-solid* to *comfortable*

- KECCAK's round function unchanged since 2008
- How many rounds?
 - 24 rounds: KECCAK/SHA-3/SHAKE
 - 12 rounds: KEYAK (CAESAR submission)
 - 5 rounds: best collision attacks [Dinur et al.] [Qiao et al.]
 - 4 rounds: best pre-image attack [Guo and Liu]


⇒ let's do hashing with 12 rounds

Better exploit parallelism

- SIMD with growing widths
 - 128, 256 and soon 512 bits
- Multiple cores

⇒ let's exploit this parallelism

KANGAROOTWELVE: parallel hashing



KangarooTwelve performance in high-end CPUs

	Short input	Long input
Intel® Core™ i5-4570 (Haswell)	4.15 c/b	1.44 c/b
Intel® Core™ i5-6500 (Skylake)	3.72 c/b	1.22 c/b
Intel® Xeon Phi™ 7250 (Knights Landing)	4.56 c/b	0.74 c/b



IACR ePrint 2016/770

<https://github.com/gvanas/KeccakCodePackage>