

Introduction to sponge-based cryptography

Part 2: Keyed modes

Joan DAEMEN

STMicroelectronics and Radboud University

SPACE 2016

Hyderabad, India, December 14, 2016


Outline

- 1 Sponge
- 2 Keyed sponge
- 3 Beyond birthday-bound security
- 4 Keyed sponge, refactored
- 5 Focus on authenticated encryption
- 6 KEYAK and KETJE

Outline

- 1 **Sponge**
- 2 Keyed sponge
- 3 Beyond birthday-bound security
- 4 Keyed sponge, refactored
- 5 Focus on authenticated encryption
- 6 KEYAK and KETJE

RADIOGATÚN [Keccak team, NIST 2nd hash workshop 2006]



- XOF: eXtensible Output Function
- Problem: **expressing security claim**
- Search for random oracle but then with inner collisions

(Early) Sponge at Dagstuhl, January 2007

Screenshot:

- Description:
 - Internal state $S = (S_A, S_G) \in \mathbb{Z}_2 \times \mathbb{Z}_2^c$ with initial value $S = (0, 0)$
 - Absorbing: for each bit p of the input:

$$S = f(S_A + p, S_G)$$
 - Resting:

$$S = f(S_A + 1, S_G)$$
 - Squeezing: for each bit z of the output:

$$z = S_A$$

$$S = f(S_A + 0, S_G)$$
- We call c : the *sponge capacity*




Generic security of Sponge [KT, Ecrypt hash, September 2007]

- Random sponges:
 - T-sponge: f is random transformation
 - P-sponge: f is random permutation
- Theorem: if no inner collisions, output is uniformly random
 - inner collision: different inputs leading to same inner state
 - Probability of inner collision:


$$\frac{M^2}{2^{c+1}} \text{ with } M : \# \text{ calls to } f$$

Promoting sponge from reference to usage (2007-2008)

- RADIOGATÚN cryptanalysis (1st & 3rd party): not promising
- NIST SHA-3 deadline approaching ...U-turn
- Sponge with *strong* permutation f : KECCAK [KT, SHA-3, 2008]




Distinguishing random sponge from random oracle



- Distinguishing advantage: $2^{-c-1}M^2$
- Problem: in real world, adversary has access to f

Differentiating random sponge from random oracle

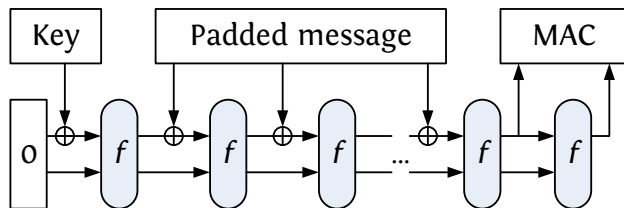


- Indifferentiability framework [Maurer, Renner & Holenstein, 2004]
- Applied to hashing [Coron, Dodis, Malinaud & Puniya, 2005]
- Random oracle augmented with *simulator* for sake of proof
- Differentiating advantage: $M^2/2^{c+1}$ [KT, Eurocrypt 2008]


Outline

- 1 Sponge
- 2 Keyed sponge**
- 3 Beyond birthday-bound security
- 4 Keyed sponge, refactored
- 5 Focus on authenticated encryption
- 6 KEYAK and KETJE

Message authentication codes




Stream encryption




- Long output stream per IV: similar to OFB mode
- Short output stream per IV: similar to counter mode

Authenticated encryption: spongeWrap [KT, SAC 2011]



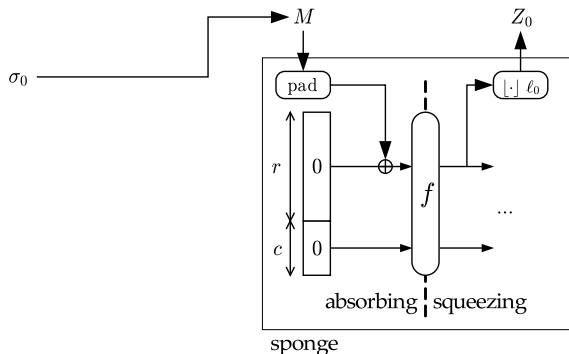
- Adopted by several CAESAR candidates
- But this is no longer sponge

The duplex construction [KT, SAC 2011]




Generic security equivalent to that of sponge

Generating duplex responses with a sponge



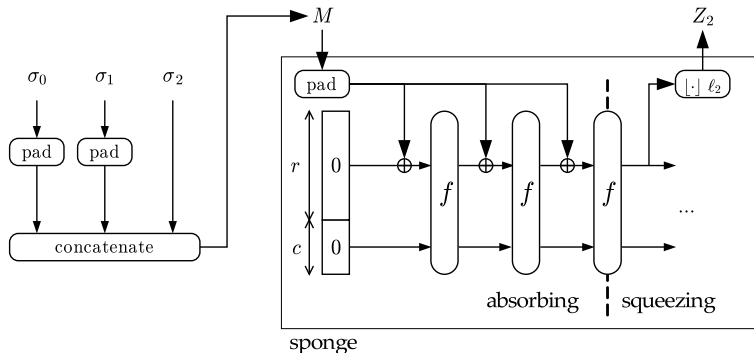
$$Z_0 = \text{sponge}(\sigma_0, \ell_0)$$

Generating duplex responses with a sponge




$$Z_1 = \text{sponge}(\text{pad}(\sigma_0) || \sigma_1, \ell_1)$$

Generating duplex responses with a sponge



$$Z_2 = \text{sponge}(\text{pad}(\sigma_0) || \text{pad}(\sigma_1) || \sigma_2, l_2)$$

Keyed sponge: distinguishing setting




- Straightforward bound: $M^2/2^{c+1} + M/2^k$
- Security strength s : expected complexity of succesful attack
 - strength s means attack complexity 2^s
 - bounds can be converted to security strength statements
- Here: $s \leq \min(c/2, k)$
 - e.g., $s = 128$ requires $c = 256$ and $k = 128$
 - $c/2$: birthday bound

Outline



- 1 Sponge
- 2 Keyed sponge
- 3 Beyond birthday-bound security**
- 4 Keyed sponge, refactored
- 5 Focus on authenticated encryption
- 6 KEYAK and KETJE

More fine-grained attack complexity




- Splitting attack complexity:
 - queries to construction: data complexity M
 - queries to f or f^{-1} : computational complexity N
- Our ambition around 2010: $M^2/2^{c+1} + NM/2^c + N/2^k$
- If we limit data complexity $M \leq 2^a \lll 2^{c/2}$:
 - $s \leq \min(c - a, k)$
 - e.g., $s = 128$ and $a = 64$ require $c = 192$ and $k = 128$

Intuition behind $NM/2^c$




- success probability per guess: $1/2^c$

Intuition behind $NM/2^c$




- $\mu \leq M$ instances with same partial r -bit input
- success probability per guess: $\mu/2^c$

Intuition behind $NM/2^c$



- $\mu \leq M$ instances with same partial r -bit input
- success probability per guess: $\mu/2^c$


Intuition behind $NM/2^c$



- $\mu \leq M$ instances with same partial r -bit input
- success probability per guess: $\mu/2^c$


An initial attempt [KT, SKEW 2011]

- bound: $M^2/2^{c+1} + NM/2^{c-1} + N/2^k$
- Problems and limitations
 - bound did not cover multi-target (key) attacks
 - proof did not convince reviewers
 - new variant (a.o. in CAESAR): inner-keyed sponge:




[Andreeva, Daemen, Mennink, Van Assche, FSE 2015]

- Inner/outer-keyed, multi-target (n), multiplicity μ
- Modular proof using Patarin's H-coefficient technique
- Bound: $M^2/2^{c+1} + \mu N/2^{c-1} + nN/2^k + \dots$




Full-state absorbing! [Mennink, Reyhanitabar and Vizár, Asiacrypt 2015]



- Absorbing on full permutation width does not degrade bounds
- We decided to use that insight in KEYAK v2
- But proven bounds had some limitations and problems:
 - term $\mu N/2^k$ rather than $\mu N/2^c$
 - no multi-key security
 - multiplicity μ only known a posteriori

Full-state absorbing! [Mennink, Reyhanitabar and Vizár, Asiacrypt 2015]




- Absorbing on full permutation width does not degrade bounds
- We decided to use that insight in KEYAK v2
- But proven bounds had some limitations and problems:
 - term $\mu N/2^k$ rather than $\mu N/2^c$
 - no multi-key security
 - multiplicity μ only known a posteriori

Outline


- 1 Sponge
- 2 Keyed sponge
- 3 Beyond birthday-bound security
- 4 Keyed sponge, refactored**
- 5 Focus on authenticated encryption
- 6 KEYAK and KETJE

The new core: (full-state) keyed duplex




- Full-state absorbing, no padding: $|\sigma| = b$
- Initial state: concatenation of key k and IV
- Multi-key: k selected from an array \mathbf{K} with index δ
- Re-phased: f, Z, σ instead of σ, f, Z
- \approx all keyed sponge functions are modes of this

Generic security of keyed duplex: the setup



- Ideal function: Ideal eXtendable Input Function (IXIF)
 - \mathcal{RO} -based object with duplex interface
 - Independent outputs Z for different paths
- Further refine adversary's capability
 - L : # queries to keyed duplex/ \mathcal{RO} with repeated path
 - q_{IV} : \max_{IV} # init queries with different keys

Generic security of keyed duplex: the bound



$$L^2/2^{c+1} + (L + 2\nu)N/2^c + q_{IV}N/2^k + \dots$$

with ν : chosen such that probability of ν -wise multi-collision in set of M r -bit values is negligible

Joint work with Gilles Van Assche and Bart Mennink, in submission

Application: counter-like stream cipher

- Only init calls, each taking Z as keystream block
- IV is nonce, so $L = 0$
- Assume $M \lll 2^{r/2}$: $\nu = 1$

Bound:

$$(2\nu)N/2^c + q_{IV}N/2^k + \dots$$

Strength:

$$s \leq \min(c - 1, k - \log_2(q_{IV}))$$

Application: lightweight MAC

- Message padded and fed via IV and σ blocks
- t -bit tag, squeezed in chunks of r bits: $c = b - r$
- adversary chooses IV so $L \approx M = 2^a$
- q_{IV} is total number of keys n

Bound:

$$M^2/2^{c+1} + MN/2^{c-1} + nN/2^k + \dots$$

Strength:

$$s \leq \min(b - a - r - 1, k - \log_2(n))$$

Imposes a minimum width of the permutation:

$$b > s + a + r$$

Outline


- 1 Sponge
- 2 Keyed sponge
- 3 Beyond birthday-bound security
- 4 Keyed sponge, refactored
- 5 Focus on authenticated encryption**
- 6 KEYAK and KETJE

What is authenticated encryption (AE)?

- Messages and cryptograms
 - $M = (AD, P)$ message with associated data and plaintext
 - $M_c = (AD, C, T)$ cryptogram with assoc. data, ciphertext and tag
- All of M is authenticated but only P is encrypted
 - wrapping: M to M_c
 - unwrapping: M_c to M
- Symmetric cryptography: same key used for both operations
- Authentication aspect
 - unwrapping includes *verification* of tag T
 - if not valid, it returns an error \perp
- Note: this is usually called AEAD

The CAESAR competition

- Public competition for AE schemes
 - consortium from academia and industry
 - aims for portfolio instead of single winner
 - CAESAR committee (secretary Dan Bernstein)
- Timeline
 - submission deadline: March 15, 2014
 - end of round 1: July 7, 2015
 - end of round 2: August 15, 2016
 - target end date: December 2017
- Status:
 - Round 1: 57 candidates
 - Round 2: 29
 - Round 3: 15 left



<http://competitions.cr.yp.to/caesar-submissions.html>

Limitations of AE

- No protection against traffic analysis
 - AE does not hide length and number of messages
 - to be addressed separately: random padding and dummy messages
- Determinism: equal messages lead to equal ciphertexts
 - information leakage
 - concern of replay attacks
 - solution: ensure message uniqueness at wrapping end
 - include nonce N in input when wrapping
 - wrapping becomes stateful
 - a simple message counter suffices
 - From now on we always include a nonce N

Functional behaviour

■ Wrapping:

- state: K and past nonces \mathcal{N}
- input: $M = (N, AD, P)$
- output: C, T or \perp
- processing:
 - if $(N \in \mathcal{N})$ return \perp
 - else add N to \mathcal{N} and return $C, T \leftarrow \text{Wrap}[K](N, AD, P)$

■ Unwrapping:


- state: K
- input: $M_c = (N, AD, C, T)$
- output: P or \perp
- processing:
 - return $\text{Unwrap}[K](N, AD, C, T)$: P if valid and \perp otherwise

Sessions

- Session: tag in cryptogram authenticates also previous messages
 - full sequence of messages since the session started
- Additional protection against:
 - insertion,
 - omission,
 - re-ordering of messages within a session
- Attention point: *last* message of session
- Alternative views:
 - split of a long cryptogram in shorter ones
 - intermediate tags

See [Bellare, Kohno and Namprempe, ACM 2003], [Keccak Team, SAC 2011], [Boldyreva, Degabriele, Paterson, Stam, EC 2012] and [Hoang, Reyhanitabar, Rogaway and Vizár, 2015]

Functional behaviour, with sessions



- Session start: creation of stateful session object D
 - if $(N \in \mathcal{N})$ (past nonces) return \perp
 - else add N to \mathcal{N} and create D with $\text{STATE} \leftarrow \text{Start}(K, N)$
- Wrapping
 - return $C^{(i)}, T^{(i)} \leftarrow D.\text{Wrap}(AD^{(i)}, P^{(i)})$
 - this updates STATE
- Unwrapping
 - return $D.\text{Unwrap}(AD^{(i)}, C^{(i)}, T^{(i)})$: $P^{(i)}$ or \perp
 - in case of no error, this updates STATE

Why (session-based) authenticated encryption?

- Convenience
 - often both confidentiality and integrity are needed
 - one scheme to choose instead of two
- Efficiency
 - combination can be more efficient than sum of the two, e.g.,
 - CBC encryption and CMAC: 2 block cipher calls per input block
 - OCB3 AE: 1 block cipher call per input block
 - sponge-based AE: 1 permutation call per input block
- Reduction of attack surface
 - differential attacks limited to session setup due to nonce
 - chosen ciphertext attacks ineffective due to \perp
- Increase of robustness against fault attacks
 - in wrap due to nonce requirement
 - in unwrap due to \perp

An ideal AE scheme

- Underlying primitive: random oracle \mathcal{RO}
 - output length ℓ implied by the context
 - $\mathcal{RO}_e(\cdot) = \mathcal{RO}(\cdot||1)$ for encryption
 - $\mathcal{RO}_a(\cdot) = \mathcal{RO}(\cdot||0)$ for tag computation
- Wrapping
 - if $(N \in \mathcal{N})$ it return \perp
 - $C \leftarrow \mathcal{RO}_e(K||N||AD) \oplus P$
 - $T \leftarrow \mathcal{RO}_a(K||N||AD||P)$
- Unwrapping
 - $P \leftarrow \mathcal{RO}_e(K||N||AD) \oplus C$
 - $T' \leftarrow \mathcal{RO}_a(K||N||AD||P)$
 - If $(T' \neq T)$ return \perp , else return P
- Note: \mathcal{RO} input shall be uniquely decodable in K, N, AD & P

Ideal AE scheme, now supporting sessions

- Starting the session
 - if $(N \in \mathcal{N})$ it return \perp
 - $\text{History} \leftarrow K || N$
- Wrapping of $M^{(i)} = (AD^{(i)}, P^{(i)})$
 - $\text{History} \leftarrow \text{History} || AD^{(i)} || 1$ and $C^{(i)} \leftarrow \mathcal{RO}(\text{History}) \oplus P^{(i)}$
 - $\text{History} \leftarrow \text{History} || P^{(i)} || 0$ and $T^{(i)} \leftarrow \mathcal{RO}(\text{History})$
 - return $(C^{(i)}, T^{(i)})$
- Unwrapping of $M_c^{(i)} = (AD^{(i)}, C^{(i)}, T^{(i)})$
 - save current state in case of error: $S' \leftarrow \text{History}$
 - $\text{History} \leftarrow \text{History} || AD^{(i)} || 1$ and $P^{(i)} \leftarrow \mathcal{RO}(\text{History}) \oplus C^{(i)}$
 - $\text{History} \leftarrow \text{History} || P^{(i)} || 0$ and $\tau \leftarrow \mathcal{RO}(\text{History})$
 - if $(\tau = T^{(i)})$ return $P^{(i)}$,
 - else $\text{History} \leftarrow S'$ and return \perp
- Note: History shall be uniquely decodable in $K, N, AD^{(i)}$ & $P^{(i)}$

Security of the ideal AE scheme

- Attack model: adversary can adaptively query:
 - Start, respecting nonce uniqueness (not counted),
 - D.Wrap (q_w times) and D.Unwrap (q_u times)
 - $\mathcal{RO}(x)$: n times
- Input to $\mathcal{RO}(K||\cdot)$ never repeats: outputs are uniformly random
 - intra-session: each input to \mathcal{RO} is longer than previous one
 - inter-session: first part of \mathcal{RO} input (N, K) never repeated
 - So ciphertexts $C^{(i)}$ and tags $T^{(i)}$ are uniformly random

Security of our ideal AE scheme (cont'd)

- Forgery:
 - building sequence of valid ciphertexts $M_c^{(1)} \dots M_c^{(\ell)}$
 - not obtained from calls to wrap for some $M^{(1)} \dots M^{(\ell)}$
- Privacy break:
 - learning on plaintext bits of $M_c^{(\ell)}$
 - without unwrapping all of $M_c^{(1)} \dots M_c^{(\ell)}$
- Complete security breakdown: key recovery
 - single target key: getting one specific key
 - multiple target: getting one key out of m target keys

Security of our ideal AE scheme (cont'd 2)

- Forgery
 - best strategy: send random but well-formatted cryptograms
 - success probability for q_u attempts: $q_u 2^{-|T|}$
- Privacy break
 - best strategy: unwrap cryptograms with modified C_i or T_i
 - success probability for q_u attempts: $q_u 2^{-|T|}$
- Key retrieval
 - best strategy: exhaustive key search
 - single target: success prob. for n key guesses $\approx n 2^{-|K|}$
 - multi-target: success prob. for n key guesses $\leq (m+1) n 2^{-|K|}$
 - Remedy against multi-target security erosion: global nonce
- Summary:
 - 1-of- m key recovery after $2^{|K| - \log_2(m+1)}$ offline calls to $\mathcal{RO}(\cdot)$
 - single privacy break/forgery after $2^{|T|}$ online calls to D.Unwrap

Instantiating our ideal AE scheme

- Replace \mathcal{RO} by full-state keyed duplex calling e.g. KECCAK- f
- Due to distinguishing bound :
 - key recovery: $\min(2^{|K|-\log_2 m}, 2^{c-\epsilon})$ offline calls to f
 - privacy break/forgery: $\min(2^{|T|}, 2^{c/2})$ online calls to f
 - ... assuming f (i.e., KECCAK- f) has no exploitable properties
- Practical scheme?
 - History includes all previous messages
 - storing it may require huge buffer
- Practical scheme!
 - keyed duplex is hard to distinguish from \mathcal{RO}
 - it compresses all History in its b -bit state S
 - at any point S : keyed hash of History
 - instantiations: our CAESAR submission KEYAK (and KETJE)

Instantiating our ideal AE scheme

- Replace \mathcal{RO} by full-state keyed duplex calling e.g. $\text{KECCAK-}f$
- Due to distinguishing bound :
 - key recovery: $\min(2^{|K|-\log_2 m}, 2^{c-\epsilon})$ offline calls to f
 - privacy break/forgery: $\min(2^{|T|}, 2^{c/2})$ online calls to f
 - ... assuming f (i.e., $\text{KECCAK-}f$) has no exploitable properties
- Practical scheme?
 - History includes all previous messages
 - storing it may require huge buffer
- Practical scheme!
 - keyed duplex is hard to distinguish from \mathcal{RO}
 - it compresses all History in its b -bit state S
 - at any point S : keyed hash of History
 - instantiations: our CAESAR submission KEYAK (and KETJE)

Advantages of sponge-based AE

- Smaller surface for cryptanalysis than block-cipher modes
 - there are no round keys
 - evolving state during session: moving target
- Cheaper protection against side channel attacks
 - DPA and DEMA limited to session setup due to nonce unicity
 - moving target during session
- Optimization of ratio security strength vs memory usage

Wish for being *online*

- Online: *being able to wrap or unwrap a message on-the-fly*
- Avoid having to buffer long messages
- Online unwrapping implies returning unverified plaintext
 - but security of our scheme relies on it
 - two ways to tackle this problem
- Tolerating Release of Unverified Plaintext (RUP)
 - catastrophic fragmentation attack [Albrecht et al., IEEE S&P 2009]
 - add security notions and attacks [Andreeva et al., ASIACRYPT 2014]
 - try to satisfy (some of) these: costly
- This can be addressed with sessions
 - split long cryptogram into short ones, each with tag
 - shorten cryptograms til they fit the unwrap buffer

Wish for surviving sloppy nonce management

- Our assumption: K, N is unique per (wrapping) Session Start
 - users/implementers do not always respect this
 - wish to limit consequences of nonce violation
- All online AE schemes leak in case of nonce violation
 - equality of first messages of session leaks in any case
 - stream encryption: re-use of keystream
 - block encryption: just equality of block(s) leaks
 - low entropy plaintexts become an issue
 - successful active attacks for quasi all proposed schemes
- Consensus among experts on following:
 - ideal security in case of nonce misuse hard to define
 - user shall be warned to not allow nonce violation
- Just avoid nonce violation


Wish for parallelism

- Many CAESAR submissions use AES
- Modern CPUs have dedicated AES instruction, e.g. AES-NI on Intel
 - pipelining: 1 cycle per round but latency of 8 to 16 cycles
 - performing a single AES: 80 cycles
 - performing 8 independent AES: 88 cycles
- Exploiting the pipeline requires ability to parallelize
- Also non-AES based schemes can benefit from parallelism, e.g.
 - pipelined architectures
 - superscalar architectures
 - SIMD instructions
- Parallelism can be supported, e.g., KEYAK

Wish for lightweight


- Whole world of buzzwords:
 - IoT, Smart Grid, RFID, ad-hoc sensor and body area network,
...
- Strongly constrained resources
 - low area: reduce chip cost
 - low power: RF powered
 - low energy: battery-life
- Specific conditions
 - short messages
 - transaction time, ...
- Compromising on
 - target security strength
 - provable security of mode
 - consequences of improper usage, ...
- Hence: **KETJE is dedicated for lightweight**

The Motorist mode of use




- SUV = Secret and Unique Value
- Plaintext absorbed in outer part, AD in inner part also
- Tag and keystream from same output block Z_i
- Specified in three layers:
 - Piston: Π of them, each one an FSKD
 - Engine: finite state machine steering the Piston(s)
 - Motorist: session starting and (un)-wrapping, using the Engine

The Motorist mode of use




- SUV = Secret and Unique Value
- Plaintext absorbed in outer part, AD in inner part also
- Tag and keystream from same output block Z_i
- Specified in three layers:
 - Piston: Π of them, each one an FSKD
 - Engine: finite state machine steering the Piston(s)
 - Motorist: session starting and (un)-wrapping, using the Engine

The Motorist mode of use




- SUV = Secret and Unique Value
- Plaintext absorbed in outer part, AD in inner part also
- Tag and keystream from same output block Z_i
- Specified in three layers:
 - Piston: Π of them, each one an FSKD
 - Engine: finite state machine steering the Piston(s)
 - Motorist: session starting and (un)-wrapping, using the Engine

The Motorist mode of use



- SUV = Secret and Unique Value
- Plaintext absorbed in outer part, AD in inner part also
- Tag and keystream from same output block Z_i
- Specified in three layers:
 - Piston: Π of them, each one an FSKD
 - Engine: finite state machine steering the Piston(s)
 - Motorist: session starting and (un)-wrapping, using the Engine

Generic security of Motorist AE session mode



Used in KEYAK v2 [KT & Ronny Van Keer, 2015]

- Plaintext absorbed in outer part, AD in inner part also
- Used in KEYAK with $c = 256$ and $b = 1600$ or $b = 800$
- Rate 544 or 1344 so we can take $\nu = 1$
- bounds:
 - nonce-respecting: $N/2^{c-1} + q_{IV}N/2^k + \dots$
 - nonce-violating: $MN/2^c + q_{IV}N/2^k + \dots$

Outline

- 1 Sponge
- 2 Keyed sponge
- 3 Beyond birthday-bound security
- 4 Keyed sponge, refactored
- 5 Focus on authenticated encryption
- 6 KEYAK and KETJE**

KEYAK [KECCAK team + Ronny Van Keer]

- AE scheme submitted to CAESAR (tweaked for round 2)
- Permutation-based mode called **Motorist**
- Makes use of KECCAK- p permutations
 - KECCAK- p : reduced-round version of KECCAK- f
 - KECCAK- f : permutations underlying KECCAK
 - all 6 functions in SHA-3 based on KECCAK- f [1600] (24 rounds)
- Generic definition with 5 parameters
 - c capacity
 - τ tag length
 - b width of KECCAK- p
 - n_r number of rounds in KECCAK- p
 - Π degree of parallelism

KEYAK named instances

- 5 named instances with $c = 256, \tau = 128, n_r = 12$
- Efficiency:
 - Short messages: Π calls to $\text{KECCAK-}p$
 - Long messages: twice as fast as SHAKE128

Name	Width b	Parallelism Π
River Keyak	800	1
Lake Keyak	1600	1
Sea Keyak	1600	2
Ocean Keyak	1600	4
Lunar Keyak	1600	8

KETJE [KECCAK team + Ronny Van Keer]

- AE scheme submitted to CAESAR (made it to round 2)
- Two instances
- Functionally similar to KEYAK
- Lightweight:
 - using reduced-round KECCAK- $f[400]$ or KECCAK- $f[200]$
 - small footprint
 - low computation for short messages
- How?
 - 96-bit or 128-bit security (incl. multi-target)
 - more ad-hoc: MONKEYDUPLEX instead of FSKD
 - reliance on nonce uniqueness for key protection

KETJE instances and lightweight features

feature		Ketje Jr	Ketje Sr
state size		25 bytes	50 bytes
block size		2 bytes	4 bytes
processing		computational cost	
session start	per session	12 rounds	12 rounds
wrapping	per block	1 round	1 round
8-byte tag comp.	per message	9 rounds	7 rounds

More on KETJE and KEYAK:
<http://ketje.noekeon.org>
<http://keyak.noekeon.org>


Safety margin of KEYAK and KETJE

S. Huang, M. Wang, X. Wang and J. Zhao, Conditional cube attack on reduced-round keccak sponge function [IACR eprint 2016/790]

- Best current cryptanalysis of keyed KECCAK- f modes
- Cube attack
 - exploits low algebraic degree of permutation
 - n rounds has degree 2^n
 - summing over inputs in affine space acts as differentiation
 - attack requires summing over around 2^n inputs
 - smart tricks allow peeling off rounds
- Most powerful attacks on KEYAK (12 rounds)
 - 7-round variant: requires 2^{42} blocks of chosen data
 - 8-round variant: requires 2^{74} blocks of chosen data

Conclusion

Permutations: good alternative for block ciphers



<http://sponge.noekeon.org/>
<http://keccak.noekeon.org/>